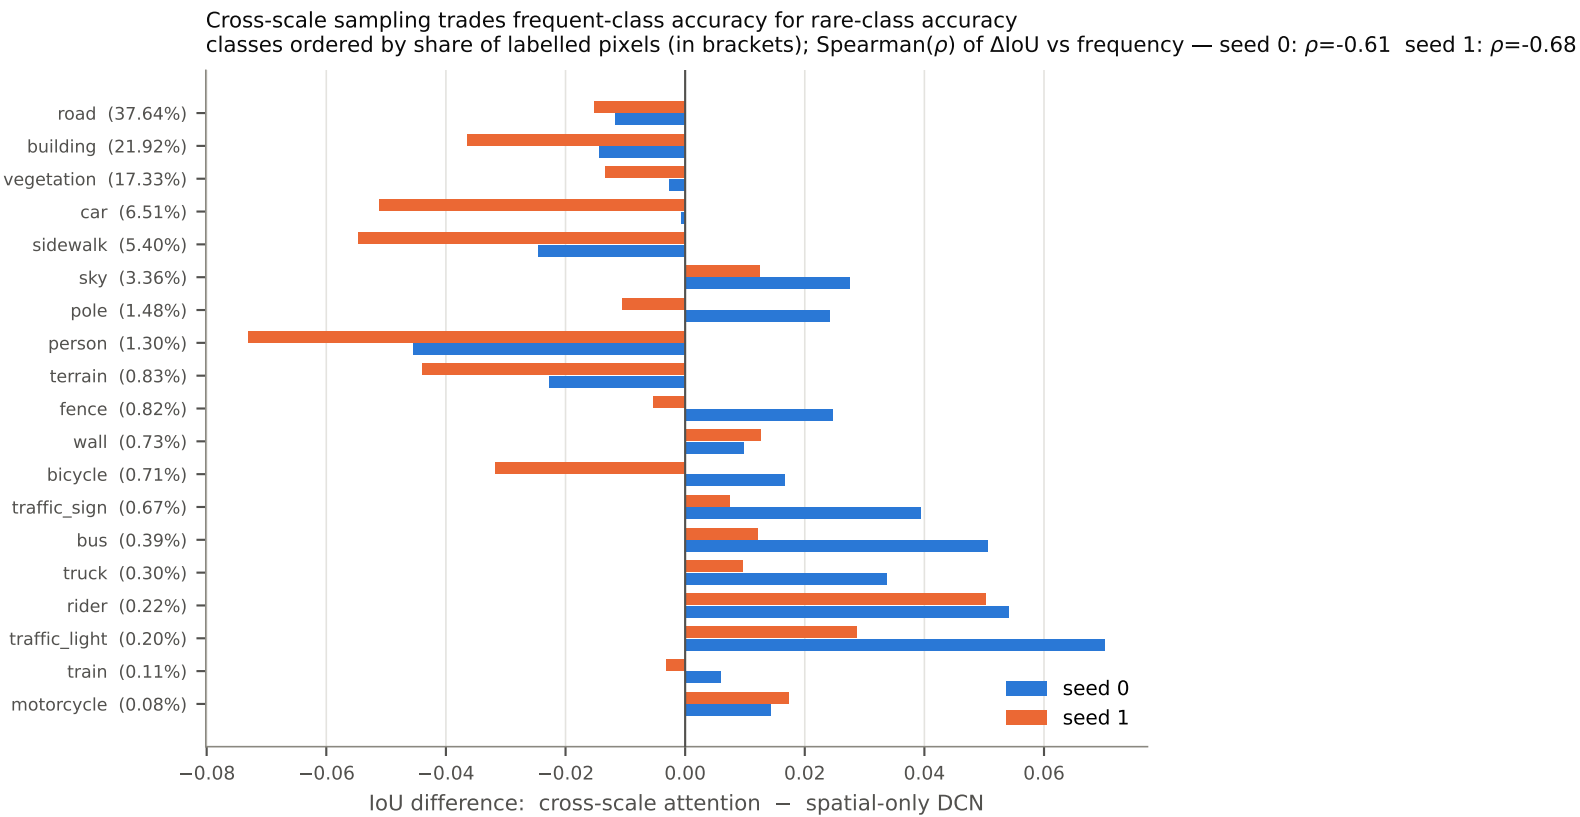
Cross-scale sampling trades frequent-class accuracy for rare-class accuracy
classes ordered by share of labelled pixels (in brackets); Spearman(ρ) of ΔIoU vs frequency — seed 0: ρ=-0.61 seed 1: ρ=-0.68
road (37.64%)
building (21.92%)
vegetation (17.33%)
car (6.51%)
sidewalk (5.40%)
sky (3.36%)
pole (1.48%)
person (1.30%)
terrain (0.83%)
fence (0.82%)
wall (0.73%)
bicycle (0.71%)
traffic_sign (0.67%)
bus (0.39%)
truck (0.30%)
rider (0.22%)
traffic_light (0.20%)
train (0.11%)
motorcycle (0.08%)
seed 0
seed 1
−0.08
−0.06
−0.04
−0.02
0.00
0.02
0.04
0.06
IoU difference: cross-scale attention − spatial-only DCN