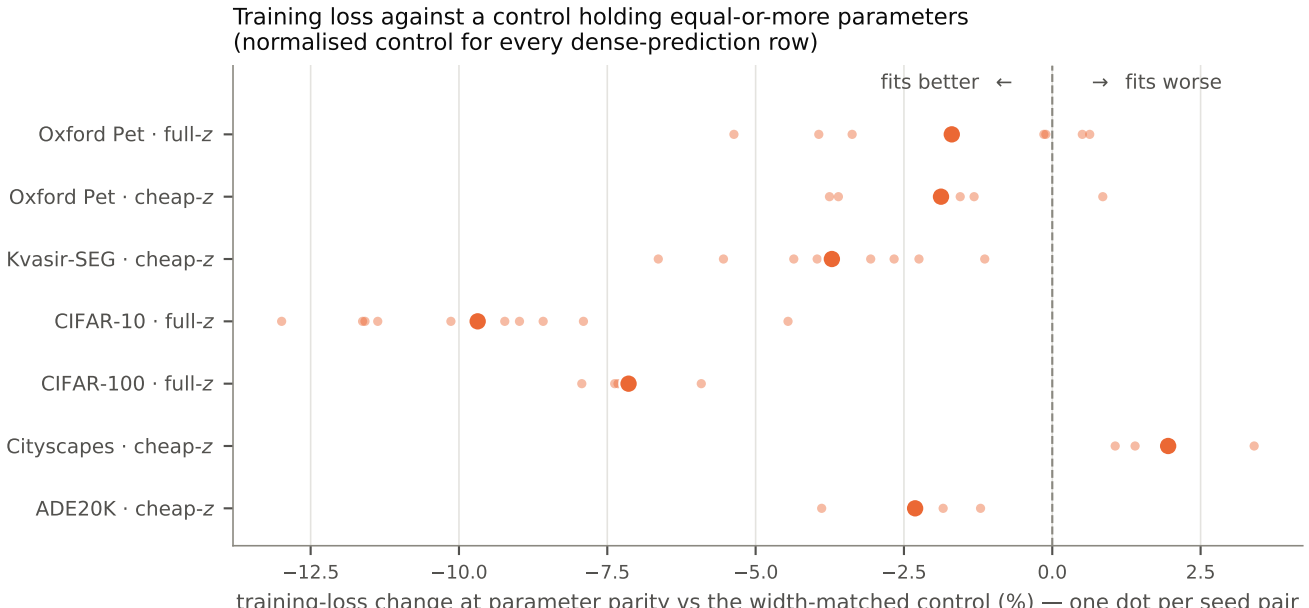
Training loss against a control holding equal-or-more parameters
(normalised control for every dense-prediction row)
fits better ←
→ fits worse
Oxford Pet · full-z
Oxford Pet · cheap-z
Kvasir-SEG · cheap-z
CIFAR-10 · full-z
CIFAR-100 · full-z
Cityscapes · cheap-z
ADE20K · cheap-z
−12.5
−10.0
−7.5
−5.0
−2.5
0.0
2.5
training-loss change at parameter parity vs the width-matched control (%) — one dot per seed pair